

Building an inverted index

Stanislav Protasov

Agenda

1. Languages
 - a. Formal approach
 - b. Tokenization
 - c. Stemming
2. Building an inverted index

Quick check

- DNS
- HTTP
- HTTPS
- Headless browser

Languages

Languages: syntax, semantics, pragmatics

- **Pragmatics:** `new_var = map(lambda x: x - 2, [4, 5, 6])`
- **Semantics:**
 - This is a valid sentence in English.
 - The worst part and clumsy looking for whoever heard light.
 - Twas brillig, and the slithy toves did gyre and gimble in the wabe.
 - Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor ...
- **Grammar (syntax):**
 - I can has cheezburger?
 - I nevr mkae tipos and erors in my sentencs.
 - I'm chuffed to bits seeing you! Do ya wanna watch some telly together, bro?
 - I'll txt w/my ETA 2U.

Syntax

Definitions

syntax guards word order

[formal] **grammar** - describes how to form strings from a language's *alphabet* that are valid according to the language's *syntax*. = set of **rules**, way to express syntax

formal language - set of all strings *allowed* by a grammar. = **satisfy** rules

Grammar is a $\langle \Sigma$ - terminals, \mathbf{N} - nonterminals, \mathbf{P} - productions, \mathbf{S} - start symbol \rangle

Chomsky Normal Form (CNF)

and Chomsky hierarchy



0. **Recursively enumerable** (any types of productions)

1. **Context-sensitive** $\alpha A \beta \rightarrow \alpha \gamma \beta$

Also noncontracting grammar ($\alpha \rightarrow \beta$, где $\alpha, \beta \in \{\Sigma \cup N\}^+$ и $|\alpha| \leq |\beta|$)

2. **Context-free** $A \rightarrow \alpha$

3. **Regular** $A \rightarrow a$ or $A \rightarrow aB \mid A \rightarrow Ba$ (exceptionally)



Extended Backus-Naur Form (EBNF)

$A = B, C.$ # concat

$A = B|C|D.$ # one of

$A = [B].$ # 0/1

$A = \{B\}.$ # 0+

$A = B\{B\}.$ # 1+

$A = (B|C)(D|E).$ # grouping (to avoid service NT)

Syntax tree

D = the | a

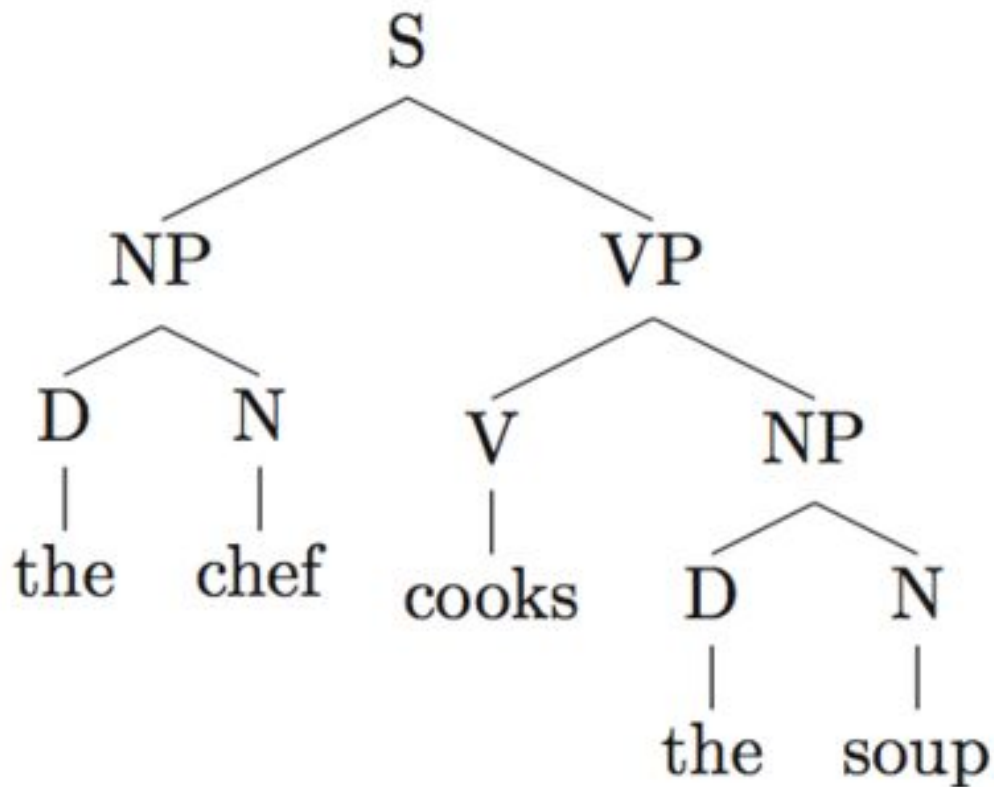
N = chef | soup

V = cooks

NP = D N | ...

VP = V N | V NP

S = NP VP



<**Σ** - terminals, **N** - nonterminals, **P** - productions, **S** - start symbol>

Tokenization

Lexemes (distinct objects of the language) are produced by **scanner**.

`token = (lexeme, token_type)`

Program, converting stream of characters into a stream of **tokens** is called **lexical analyzer, lexer, tokenizer**.

```
i like to read science fiction.  
[('i', 'PRP'), ('like', 'VBP'), ('to', 'TO'),  
 ('read', 'VB'), ('science', 'NN'), ('fiction'  
 , 'NN'), ('.', '.')] ]
```

Syntax analysis helps tokenization

L'ensemble □ one token or two? **L** ? **L'** ? **Le** ?

莎拉波娃现在居住在美国东南部的佛罗里达

استقلت الجزائر في سنة 1962 بعد 132 عام من الاحتلال الفرنسي.

language-specific punctuation

```
import nltk
```

```
nltk.download()
```

```
st = nltk.data.load('tokenizers/punkt/english.pickle')
```

```
st.tokenize(text)      # sentence splitting
```

```
nltk.download('punkt')
```

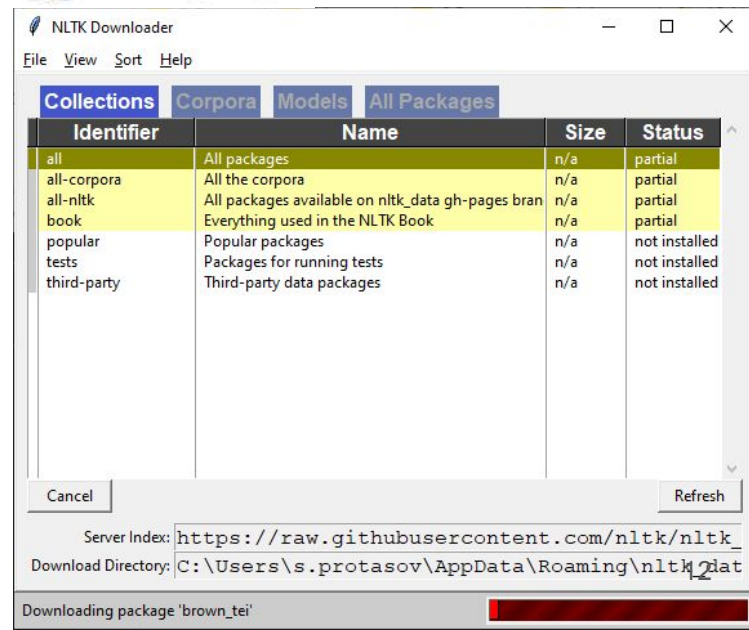
```
nltk.word_tokenize()  # language specific
```

grammar based tokenization

```
simple_grammar = nltk.parse_cfg(...)
```

```
parser = nltk.ChartParser(simple_grammar)
```

```
trees = parser.nbest_parse("A car has a door")
```



Techniques used across methods (semantics)

- **case folding:** London = london; Лев = лев
- **stemmer vs lemmer:**
 - stemming:* compress = compression = uncompressed
бегу = бег
 - lemmatization:* better = good
бегу = бегать
- ignore **stop words:** to, the, it, be, or, ...
 - Problems arise when search on “To be or not to be” or “the month of May”
- **Thesaurus:** fast = rapid; лев = лёвушка
 - handbuilt clustering

Stop here!

1. Language, Syntax, Grammar
2. Formal languages, grammars and types
3. Tokenization and syntax analysis
4. Stems and stemming

Boolean retrieval

BIR is based on **Boolean logic** and classical **set theory** in that both the documents to be searched and the user's query are conceived as sets of terms (a **bag-of-words model**). Retrieval is based on **whether or not the documents contain the query terms**.

(Wikipedia), [[The Book](#)]

- **Boolean query**

- E.g., (“obama” AND “healthcare” AND NOT “news”)

Search and recommender systems idea

Boolean retrieval, exact nearest neighbour search or exact range queries can be too **expensive**. *Can we do faster?*

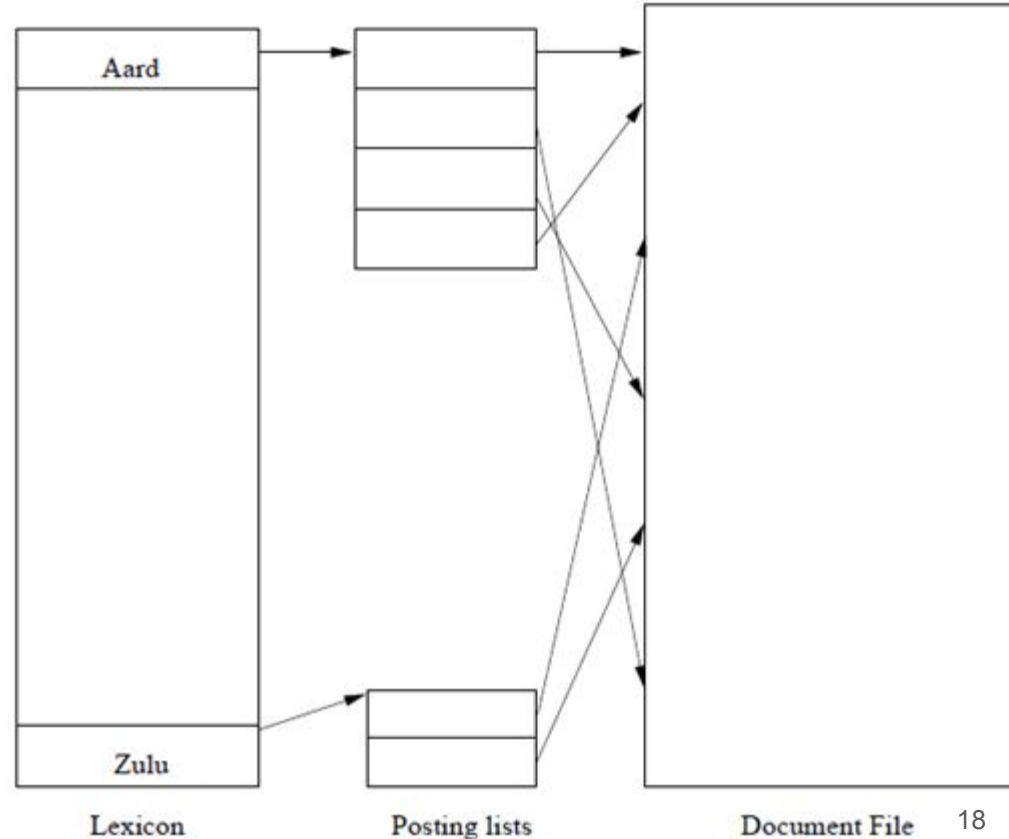
Pre-select (*pre-ranking sets*, approximate NN ...) which is done fastly (e.g. $O(\log(N))$) selects enough to catch [almost] all relevant elements. Requires data structures: indices.

Select (*ranking*, exact match) is done on a smaller set.

Text indexing: inverted index

Inverted index

- 1) Build a **lexicon** for the whole database
- 2) For each word of lexicon build a **posting list** (set of pointers)
- 3) [optional] persist this structure as a sparse matrix




On lexicon

- Languages have lots of names and phrasal forms of 2+ words (e.g. New York) → use **bigrams** and alternative tokenizations (e.g. subword tokenization)
- Stopwords should be selected carefully

Document frequency

Idea: a term is more discriminative if it occurs only in fewer documents

word frequency and rank in *Romeo and Juliet* (linear-linear)



<https://medium.com/datadriveninvestor/zipfs-law-breakdown-application-in-app-development-5e9cda70cdc8>




Figure 2.1. A plot of the hyperbolic curve relating f , the frequency of occurrence and r , the rank order (Adapted from Schultz** page 120)

On dictionary construction


External memory dictionaries are useless for real-time applications. Files used for persistence (mmap)

Index **building** require more memory then index itself. Map-Reduce is widely used for this. (Book, 4.4)

Indices are usually **static**. For dynamic read (Book, 4.5)

What if total number of documents is big for a single machine?

- Use Trie
- Use Sharding: *hash* based, *lexicographical*




Search with Boolean query

- **Boolean query**
 - E.g., “obama” AND “healthcare” AND NOT “news”
- **Pre-ranking set**
 - Lookup query term in the dictionary
 - Retrieve the posting lists
 - Operation
 - AND: intersect the posting lists ([skip-lists can help to intersect in \$O\(m+n\)\$](#))
 - OR: union the posting list
 - NOT: diff the posting list
- **Last step**
 - *Re-check* selected documents hold **expected substring** (for query search)

Deficiency of Boolean model

- The query is unlikely precise
 - “**Over-constrained**” query (terms are too specific): no relevant documents can be found
 - “**Under-constrained**” query (terms are too general): over delivery
 - It is hard to find the right position between these two extremes (hard for users to specify constraints)
- Even if it is accurate
 - Not all users would like to use such queries
 - All relevant documents are **not equally** important
 - No one would go through all the matched results
- Relevance is a matter of degree!

Document Selection vs. Ranking



Ranking is often preferred

- Relevance is a matter of degree
 - Easier for users to **find** appropriate **queries**
- A user can stop browsing anywhere, so the **boundary** is **controlled by the user**
 - Users prefer coverage would view more items
 - Users prefer precision would view only a few
- Theoretical justification: Probability Ranking Principle

Retrieval procedure in modern IR

- Boolean model provides all the ranking candidates
 - Locate documents satisfying (somehow) Boolean condition
 - E.g., “obama healthcare” -> “obama” **OR** “healthcare”
- Rank candidates by **relevance**
 - Important: the notation of relevance
- Efficiency consideration
 - Top-k retrieval ([Google](#))

Intuitive understanding of relevance

	information	retrieval	retrieved	is	helpful	for	you	everyone
Doc1	1	1	0	1	1	1	0	1
Doc2	1	0	1	1	1	1	1	0
Query	1	1	0	0	0	0	0	0



*E.g., 0/1 for Boolean models,
probabilities for probabilistic models*

Ranking over Inverted Index: TF-IDF

Term frequency $\text{tf}(t, d)$, count of a term t in a document d .

- Boolean "frequencies": $\text{tf}(t, d) = 1$ if t occurs in d and 0 otherwise;
- **term frequency adjusted** for document length : $f_{t,d} \div (\text{number of words in } d)$
- **logarithmically scaled** frequency: $\text{tf}(t, d) = \log(1 + f_{t,d})$
- augmented frequency

$$\text{tf}(t, d) = 0.5 + 0.5 \cdot \frac{f_{t,d}}{\max\{f_{t',d} : t' \in d\}}$$

Inverse document frequency $\text{idf}(t, D) = \log \frac{N}{|\{d \in D : t \in d\}|}$

$$\text{tfidf}(t, d, D) = \text{tf}(t, d) \cdot \text{idf}(t, D)$$

Read at home (extra reading)

The book, chapter 4 (index construction)