Sound and speech retrieval

Stanislav Protasov

Agenda

- Sound as a wave
- Music search
- Speech recognition

What is the sound?

Sound is a **vibration** that propagates through a transmission medium such as a gas, liquid or solid.

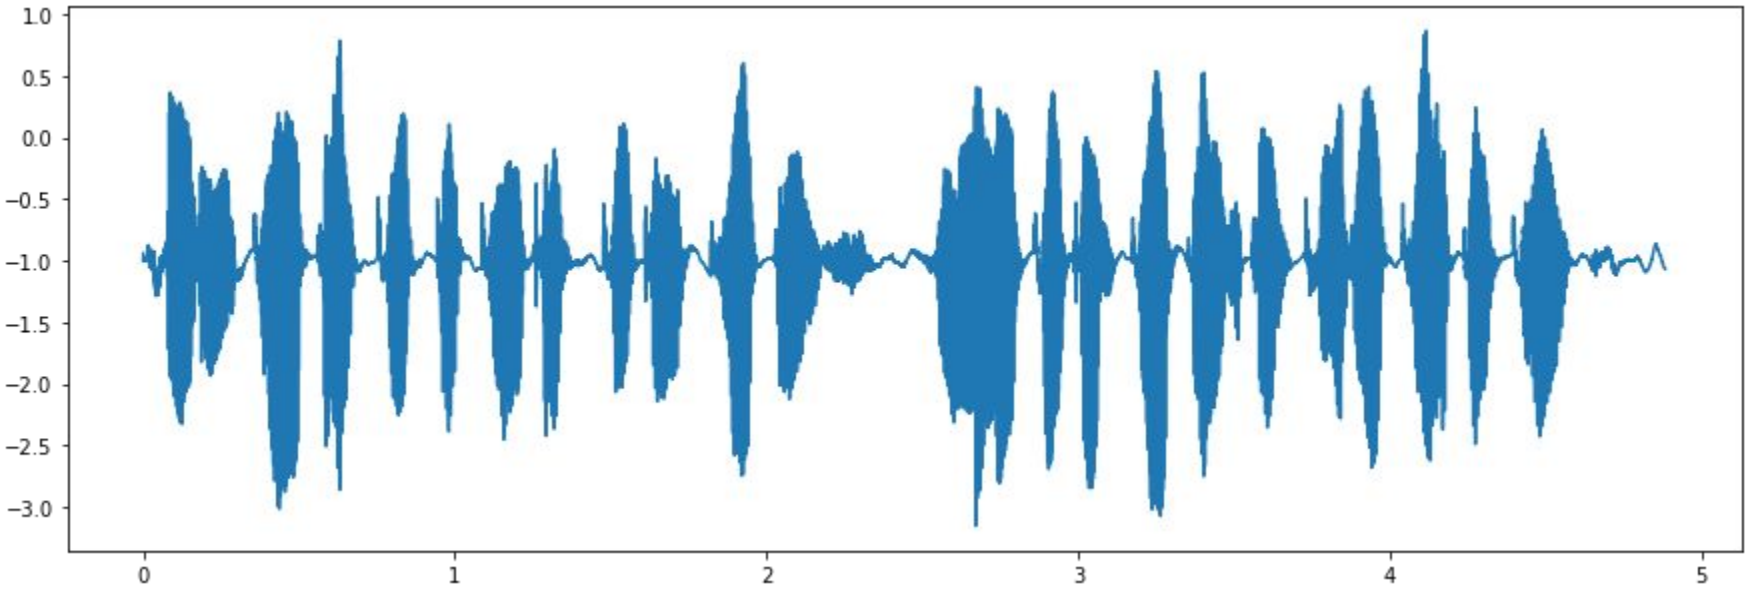

Vinyl player

National Institute on Deafness and Other Communication Disorders

How ear works

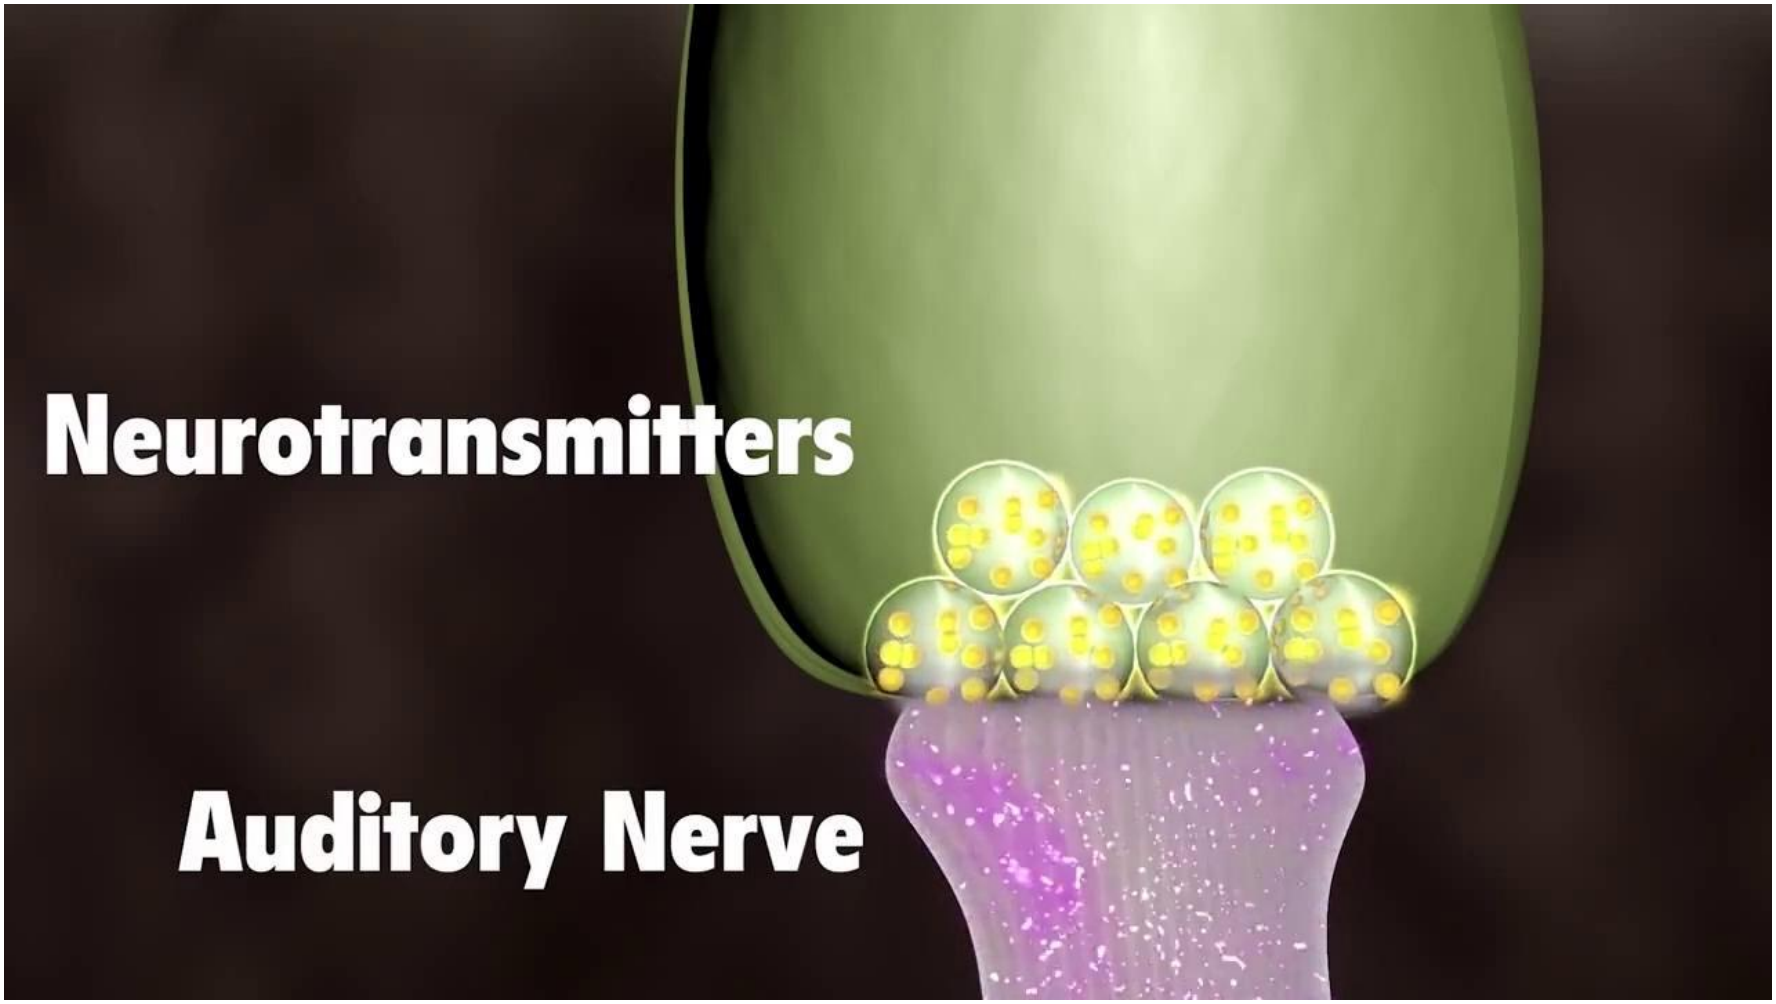

Euler's identity to link complex exponent with frequencies

$$e^{ix} = \cos x + i \sin x,$$

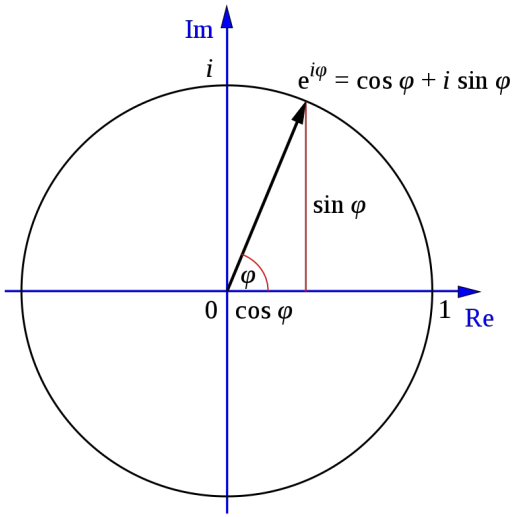

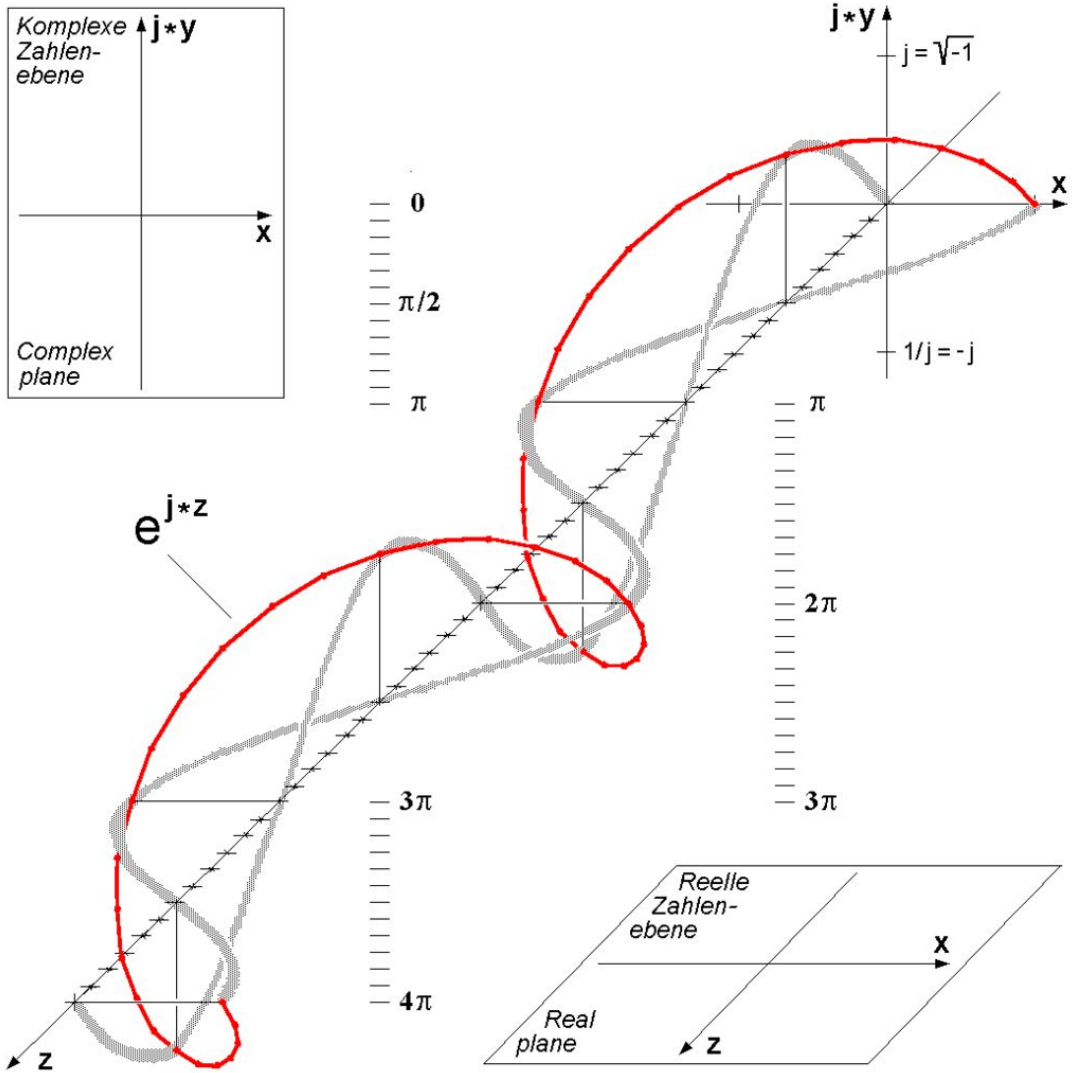

Fourier Transforms

$$FT: \quad \hat{f}(w) = \int_{-\infty}^{+\infty} f(x) e^{-2\pi i x w} dx$$

$$DTFT: \quad X_T(w) = \sum_{n=-\infty}^{+\infty} f(nT) e^{-2\pi i w n T}$$

$$DTFT + window: \quad X_T(w) = \sum_{n=0}^{M} f(nT) W\left(\frac{n}{M}\right) e^{-2\pi i w n T}$$

$$DFT: \quad X_{T,N}(k) = X_T\left(\frac{k}{NT}\right) = \qquad k = 0, 1, \ldots, N-1$$

$$= \sum_{n=0}^{M} f(nT) W\left(\frac{n}{M}\right) e^{-2\pi i \frac{kn}{N}}$$

Fourier transform

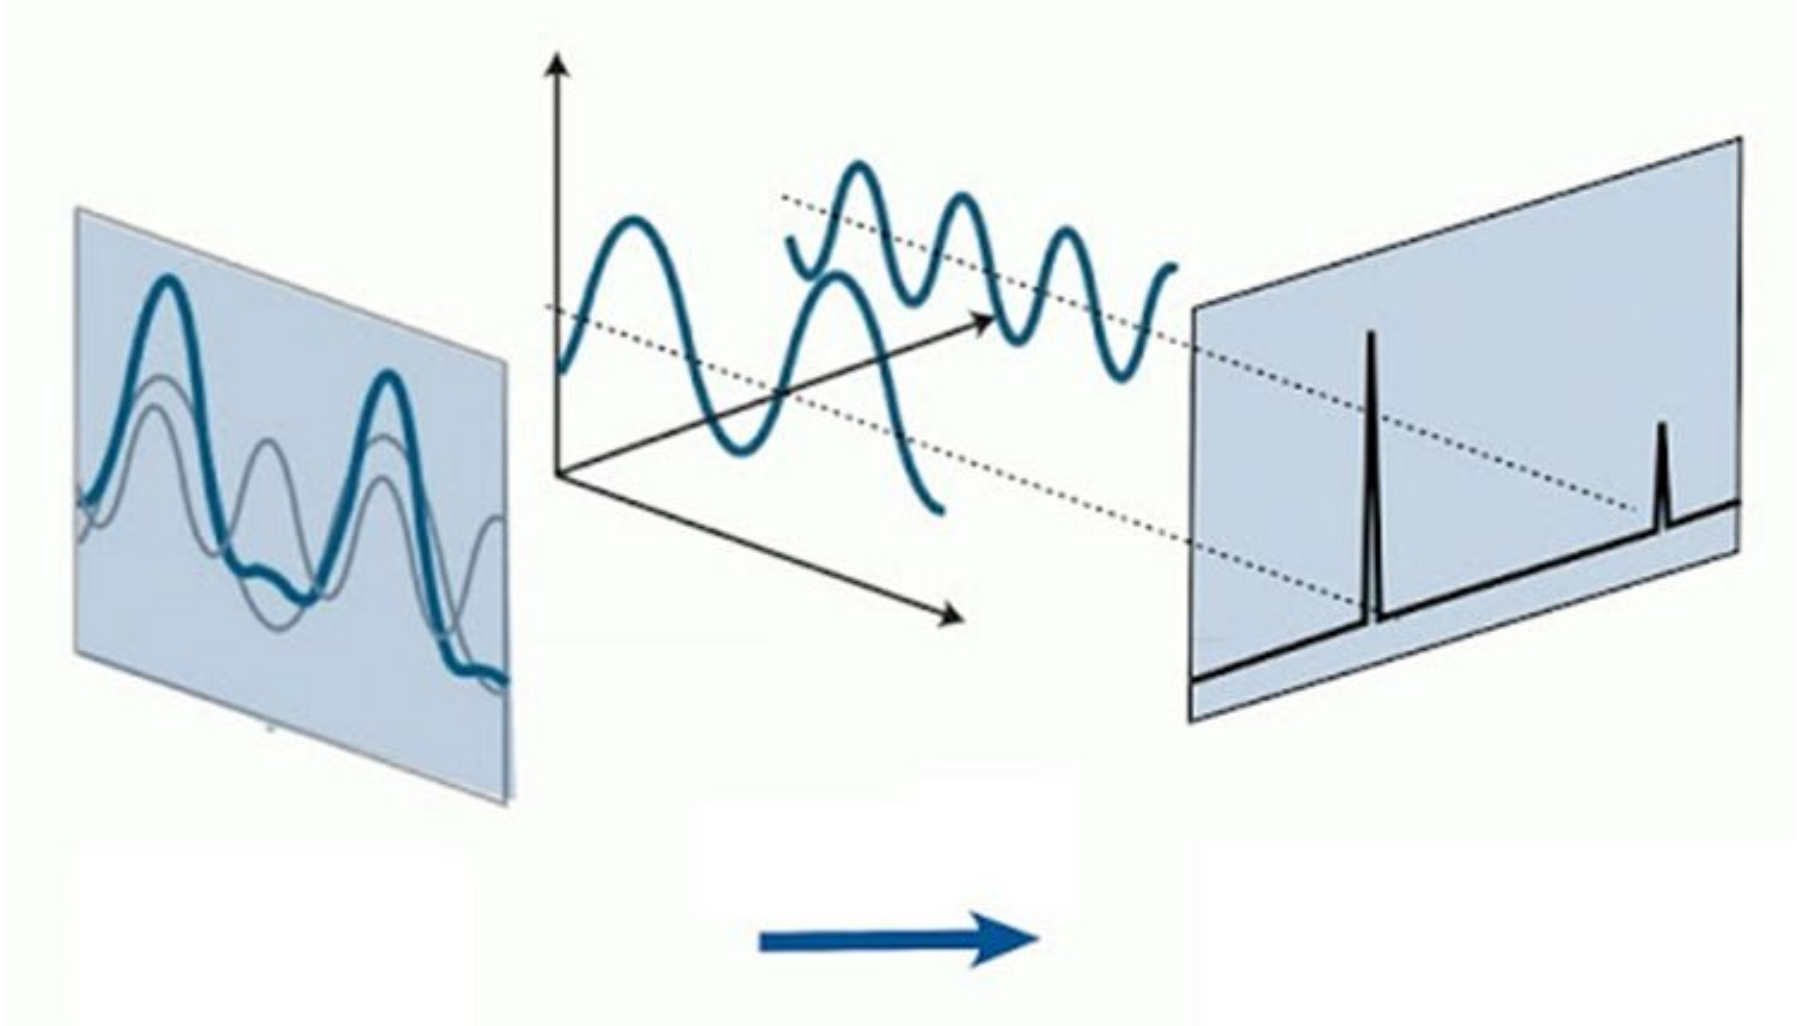

Guitar pitch

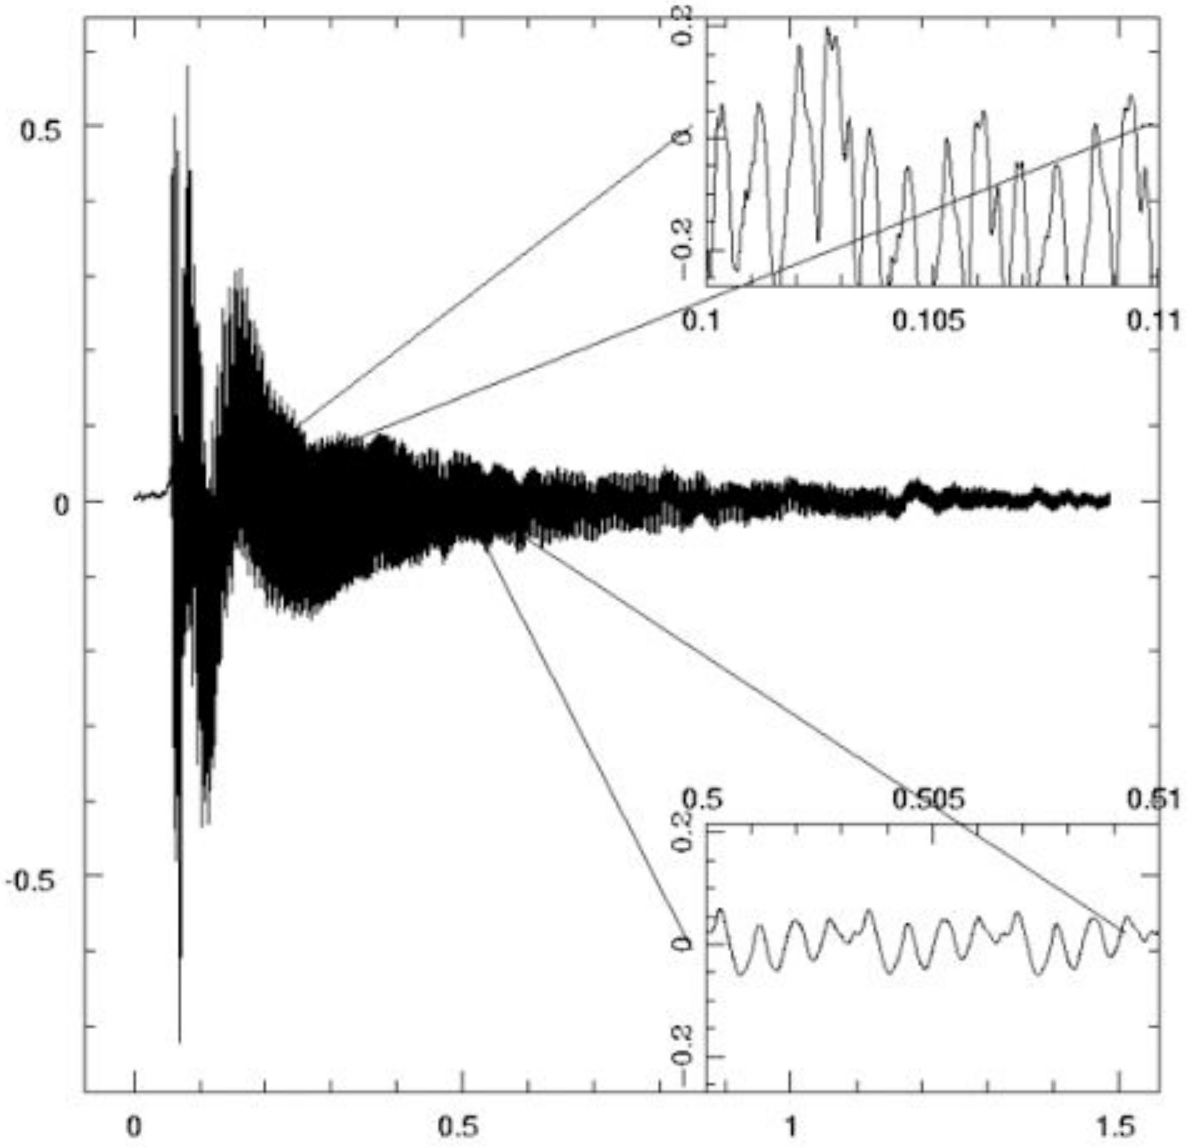

Wavelets

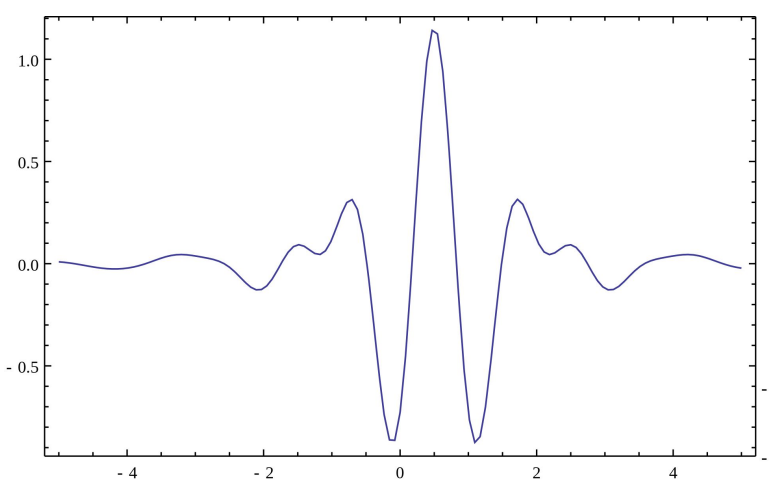
1.0
0.5
0.0
- 0.5
- 4
- 2
0
2
4

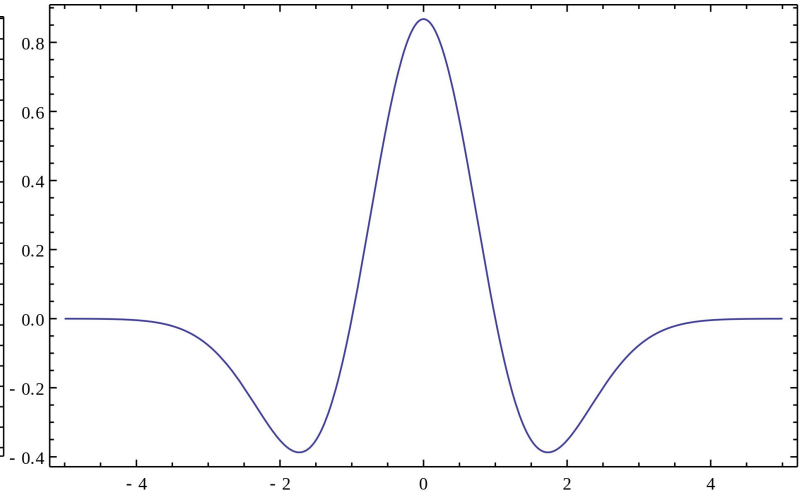
0.8
0.6
0.4
0.2
0.0
- 0.2
- 0.4
- 4
- 2
0
2
4

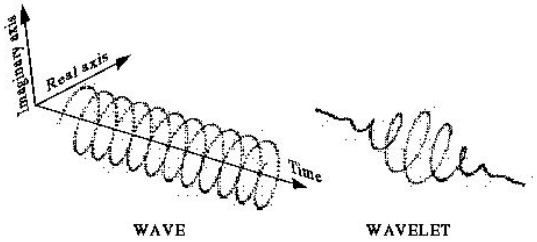
Imaginary axis
Real axis
Time
WAVE
WAVELET

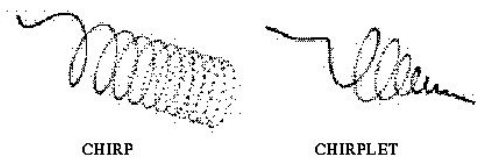
CHIRP
CHIRPLET

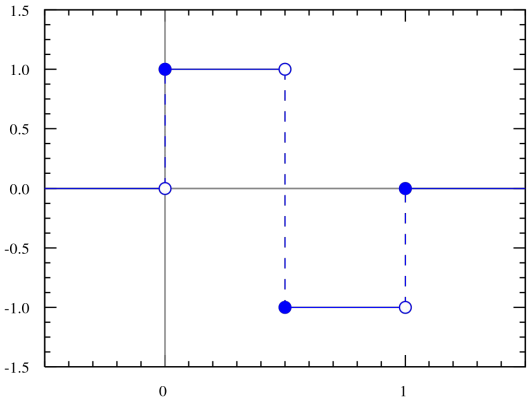
1.5
1.0
0.5
0.0
-0.5
-1.0
-1.5
0
1

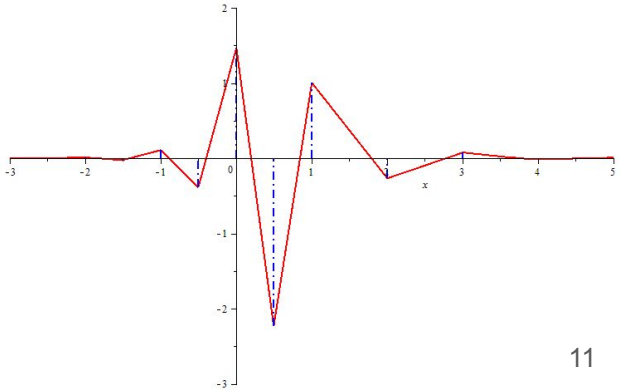
2
1
0
-1
-2
-3
-3
-2
-1
0
1
3
4
5
x

What is the sound for human?

We percept sound using **frequency** receptors. Each moment looks like this:

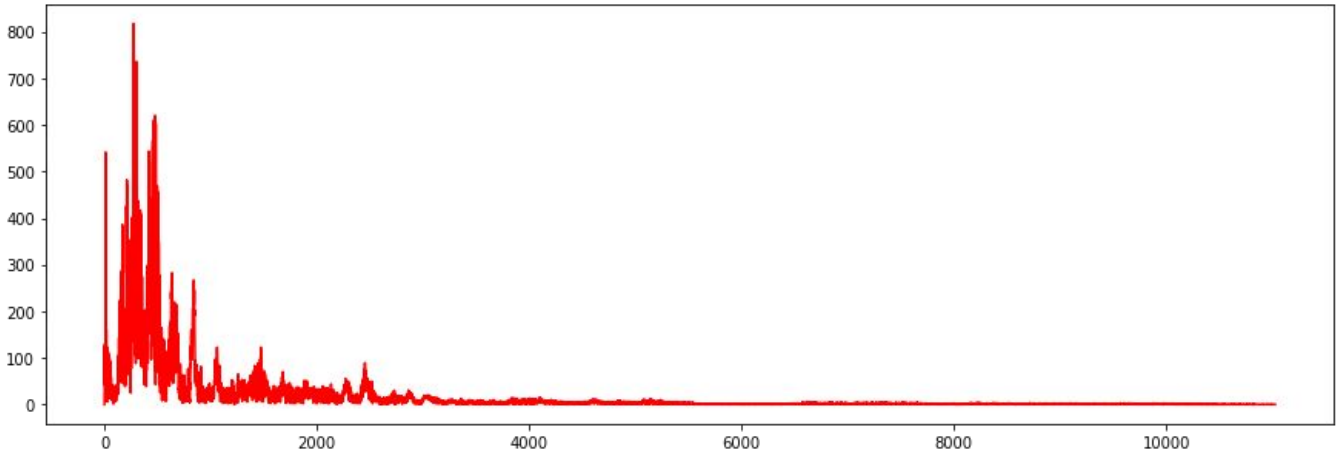

Also important — we perceive sounds in **log scale**

Timeline is like this:

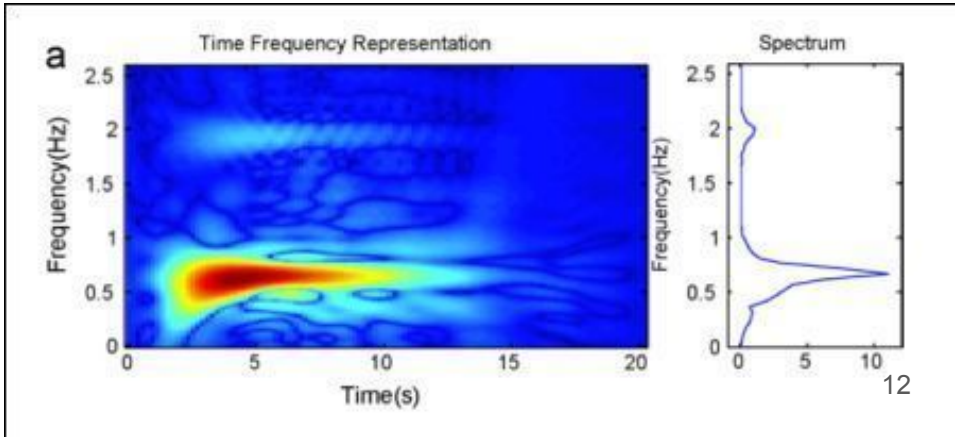

Sound *recording* and *playback*

- Digital uncompressed sound consists of regular measurements of signal.
- Measurement frequency is managed using RATE parameter
  - 22050 means 22050 measurements per second (**discretization**)
- How accurate we measure in managed is tuned with format (**quantization**)
  - How many different amplitude values can be encoded
- Channels — number of inputs/outputs (stereo=2, mono=1)
- BPS = RATE * CHANNELS * FORMAT
- Together this is **PCM — pulse code modulation**

Nyquist-Shannon (Kotelnikov) theorem

If a function **x(t)** contains no frequencies higher than ***B*** hertz, it is **completely determined** by giving its values at a series of points spaced **1/(2B)** seconds apart.

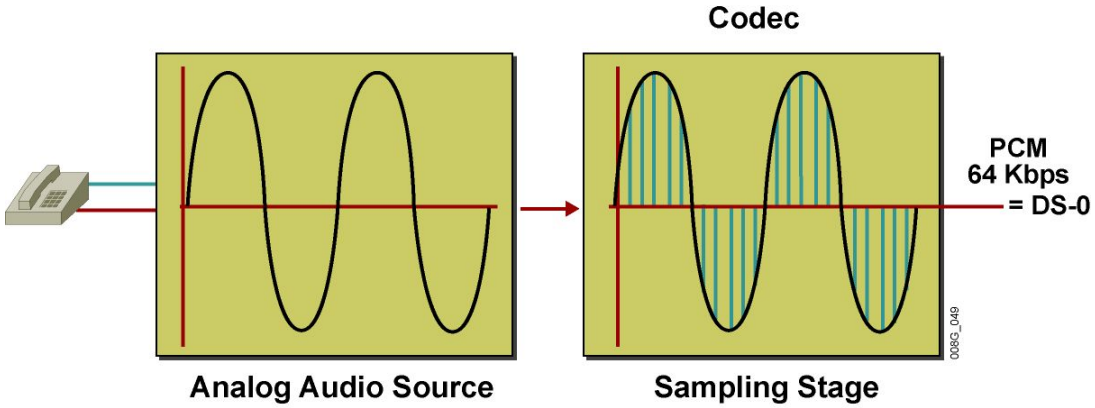

What **if contains**? Aliasing. ***n(k)***?

$$\{\sin(k\,x) = \sin(n\,x),\ n < k\}$$

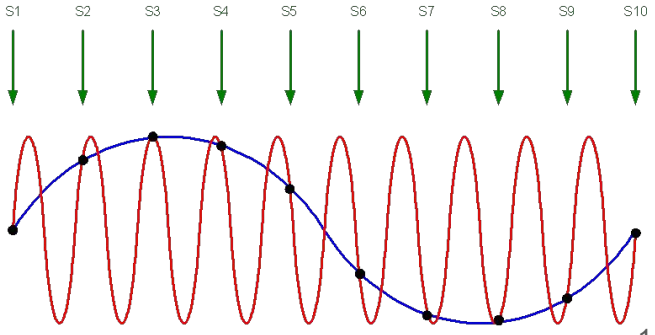

- `sin(a)+sin(b) = 2·sin(½(a+b))·cos(½(a-b))`

Music fingerprinting

Why?

- I like this song, I want to buy it
- Forensic (when was this song playing)
- **Copyrights** (see youtube or instagram policy)

How to form a query

1. Exact sample
2. Humming

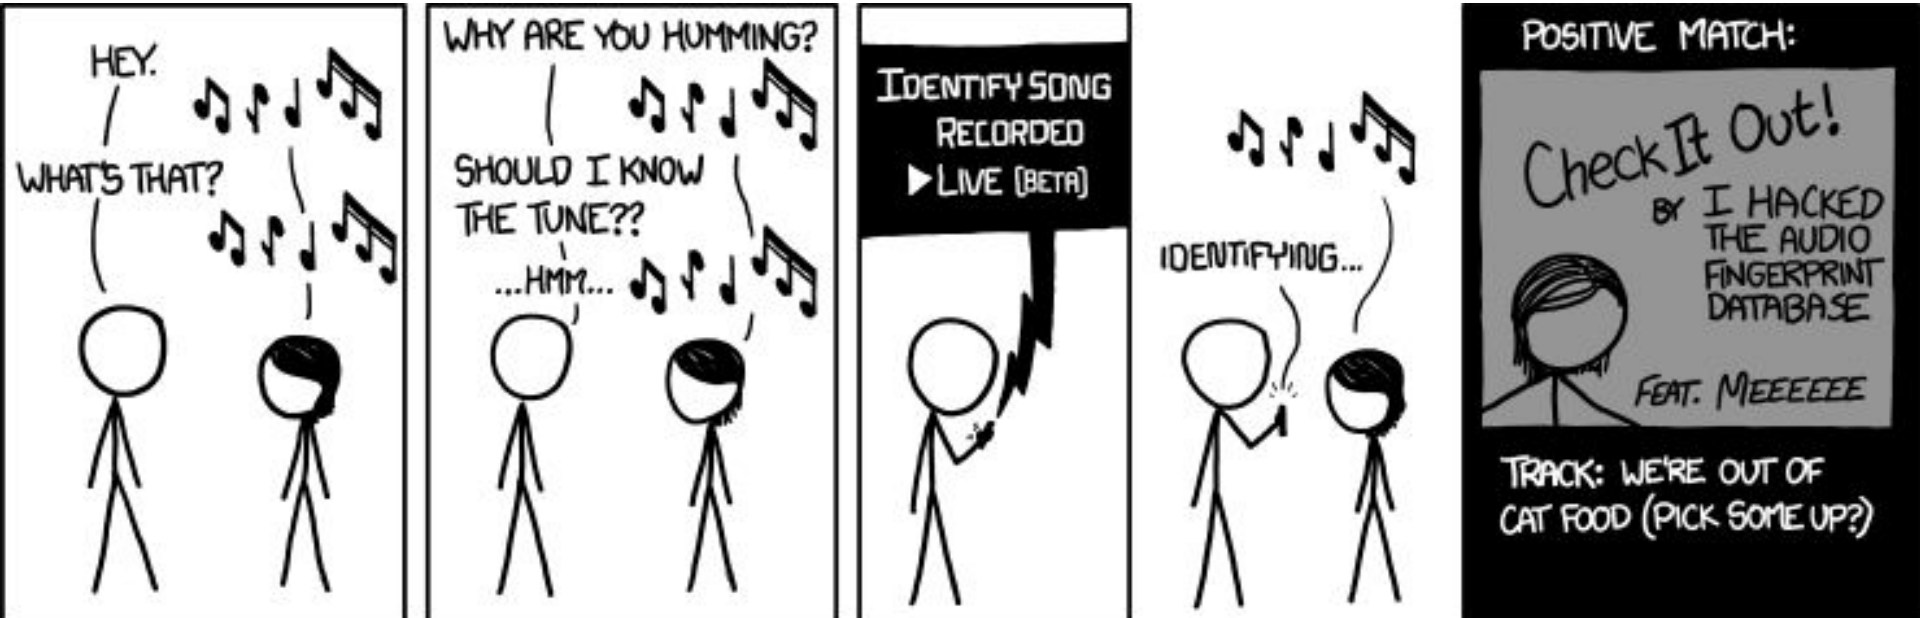

Shazam exact match algorithm (1)

1. Robust spectrogram. (Log-scale bins of frequencies)

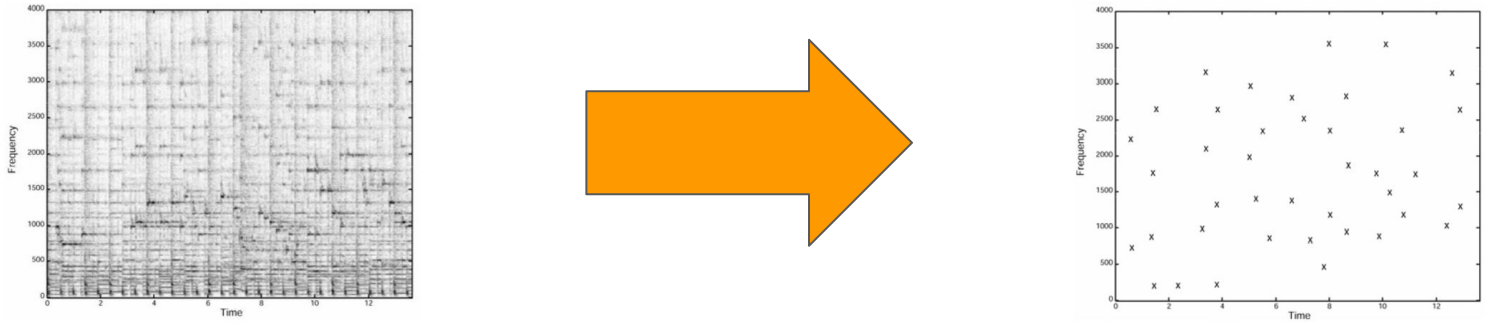

2. Build **pairs for hashing (32bit)**: anchor point + other point from target zone.

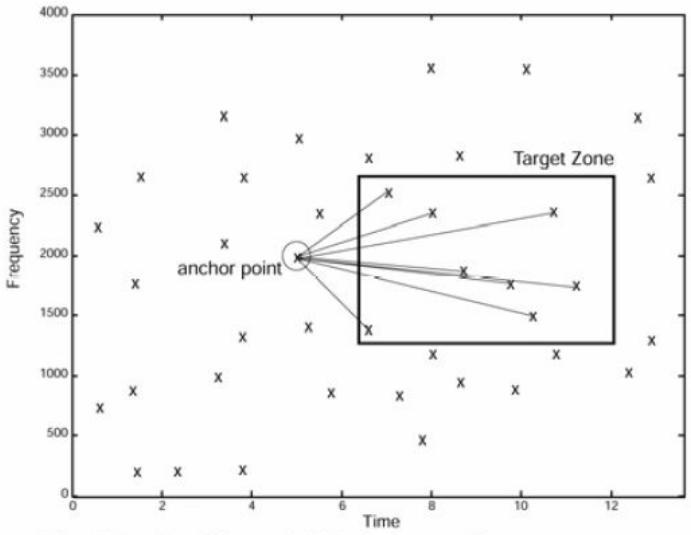

Fig. 1C - Combinatorial Hash Generation

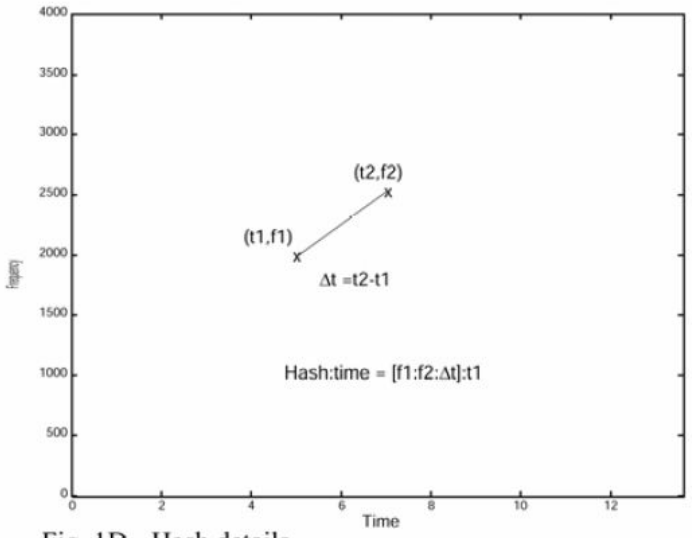

Fig. 1D - Hash details

Shazam algorithm (2)

3. Put those points to a **hashmap**.
Memory: `[4B (hash) + 4B (val)] * peaks`.

4. Query for songs with a processed sample.

5. Plot a histogram of offsets (get times from HT, put them in bins), identify real offset

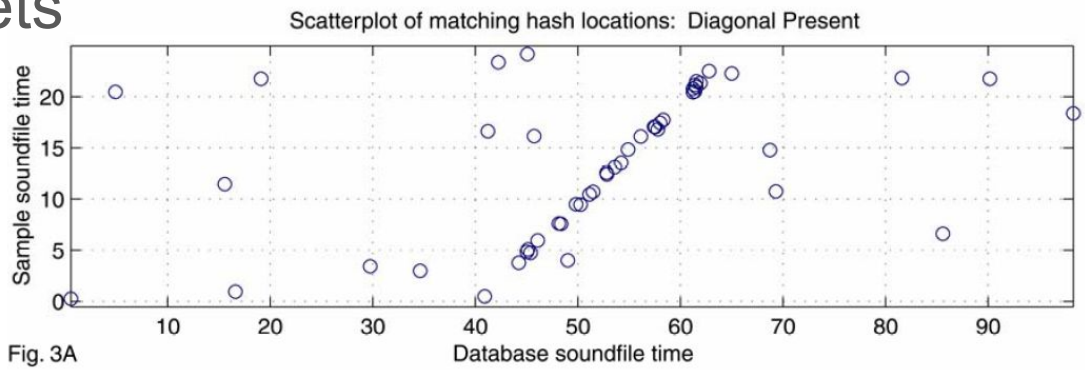

Fig. 3A

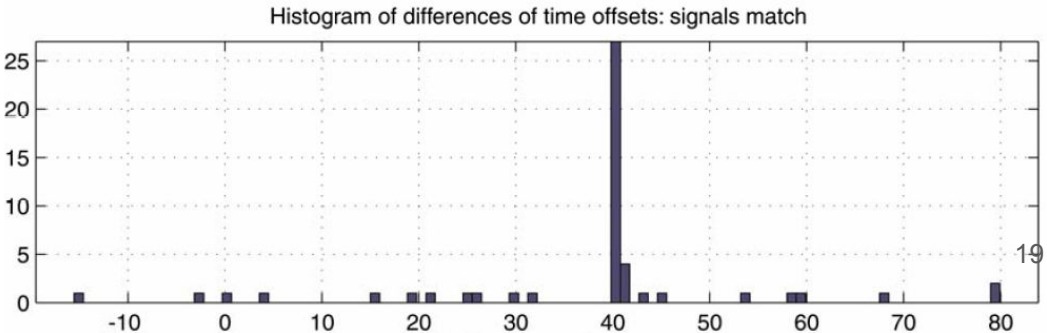

Other fingerprinting approaches

Exact frequency and tempo are not always important:

- Zero-crossing rate
- Spectrum
- Envelope (spectral flatness, frequency band)
- ...

Query by Humming (QbH)

- Detect **coarse melodic contour**, retrieve by string search
  - S=same note, U=up, D=down
  - E.g., Beethoven's 5th: – **S S D U S S D**
  - OR U/D/S – but with five contour levels

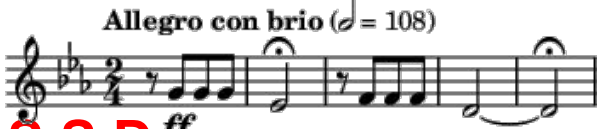

- Add **rhythm information**
- Use **beat information**
- Use **HMMs** to represent song database
- ***Dynamic Time Warping*** (DTW) based algorithm, match waveform directly

Delta pitch	2	2	1	2	-2	-1	-2	-2
IOI	1	1	1	1	1	1	1	1
IOI ratio	1	1	1	1	1	1	1	1
State	α	α	β	α	χ	δ	χ	χ

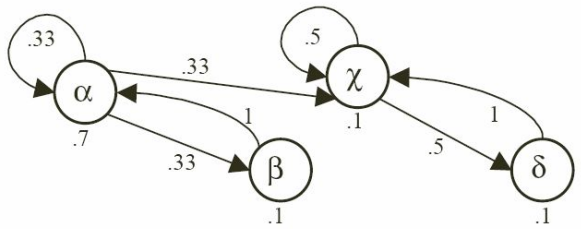

Dynamic Time Warping

- Take to sequences of lengths N and M
- Build `N*M` matrix ***d*** of distances (diffs)
- Build a matrix of deformation,

$$D_{i\ j} = d_{i\ j} + min(D_{i-1\ j}, D_{i-1\ j-1}, D_{i\ j-1}). \qquad (3)$$

- Search for a path `(1,1)-(N,M)` with minimal average value weight.

$$DTW(Q,C) = min\left\{ \frac{\sum_{k=1}^{K} d(w_k)}{K} \right\}. \qquad (4)$$

Can give false positives

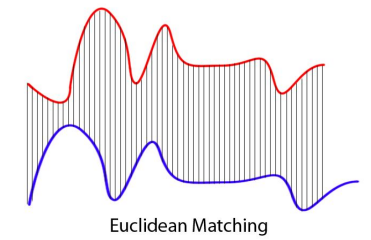

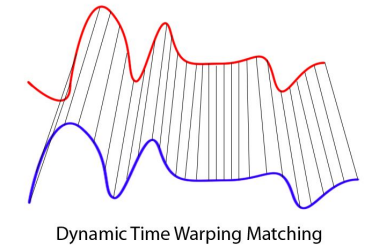

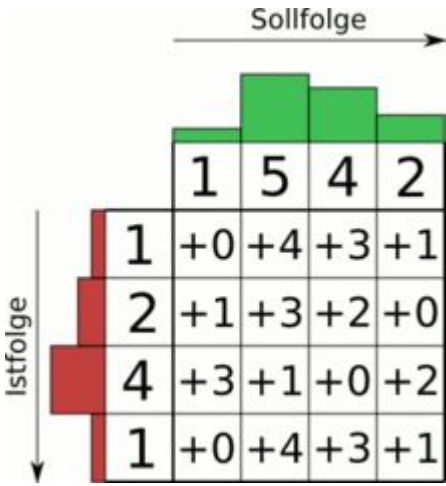

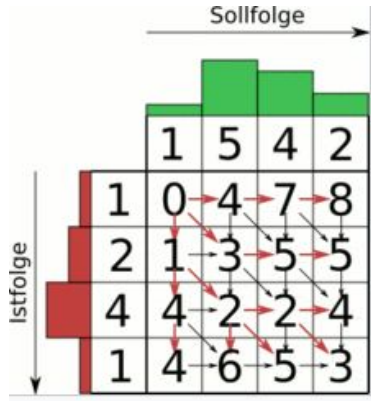

Google Hum to search

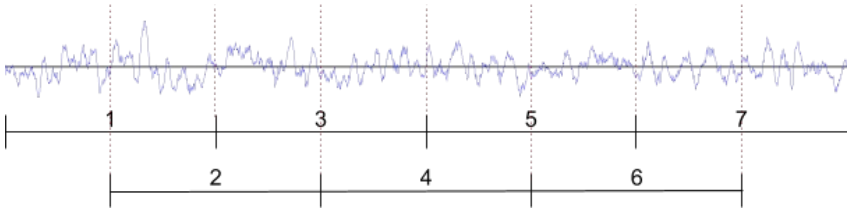

- Convolutional net to build a fingerprint-based from 8 seconds with 0.5 sec. step
- Inaccurate vector search with space partitioning and vector quantization
- Retrieve all candidate's fingerprints
  - Accurate match on all set of embeddings for candidates

Speech (text) processing

Acoustic model

As text consist of letters, speech consists of phonemes.

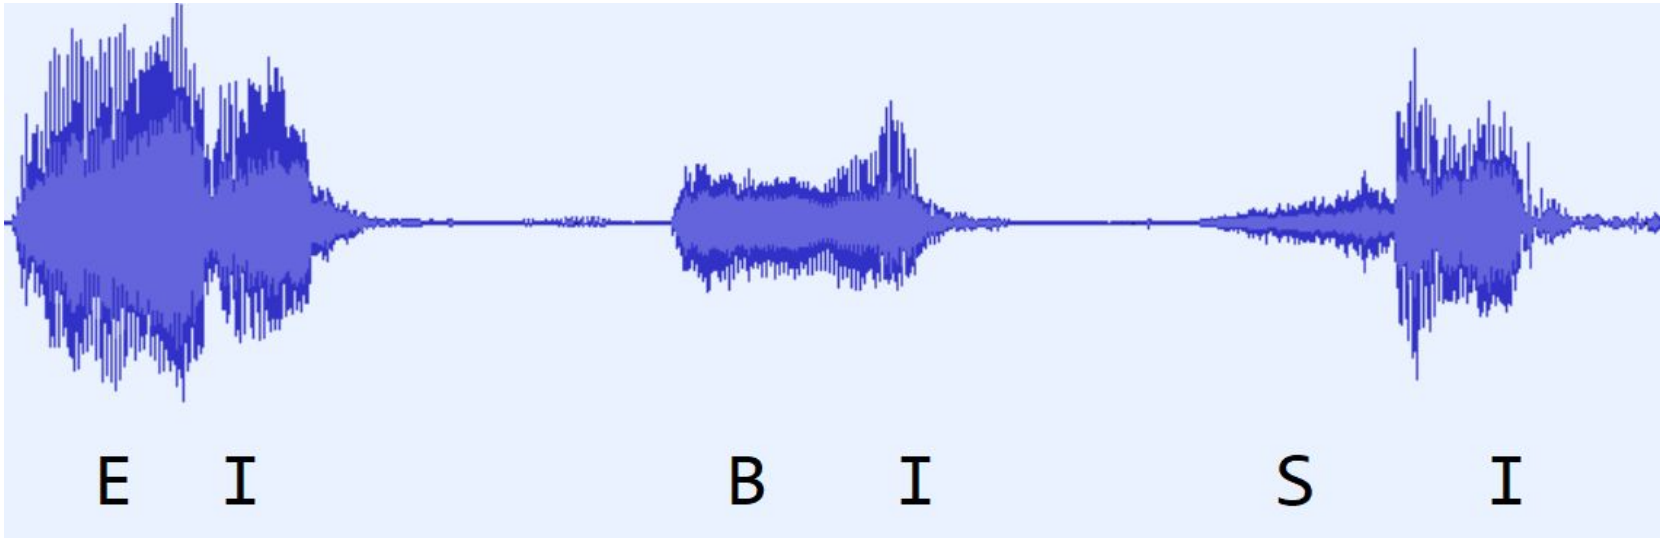

AM: spectrum → phoneme

Language model (in recognition)

Probabilistic model that predicts probability of a word given a sequence of phonemes.

Similar model is used to model sentences of words.

Speech generation

1) Text preprocessing
   a) Number to text
   b) Abbreviations to text
   c) Typo fix
2) Split text into phrases (punctuation, constructions)
3) Phonetic construction (language model)
   a) queue - [kju]
   b) Арбалетчиков
      i) a0 r b a0 lj e**1** t ch i0 k o0 v

Speech generation

1) **Accents** are set
    a) Using a dictionary
    b) Using rules
    c) Using statistics (speaker examples)
2) **Reversed acoustic model** is used to consider surrounding
3) **Timbre** is generation with **vocoder**
    a) or RNNs